

Machine Learning Foundations

Linear Algebra II: Matrix Operations

Use Tensors in Python to
Solve Systems of Equations and
Identify Meaningful Patterns in Data

Jon Krohn, Ph.D.

jonkrohn.com/talks

github.com/jonkrohn/ML-foundations



Machine Learning Foundations

Linear Algebra II: Matrix Operations

Slides: jonkrohn.com/talks

Code: github.com/jonkrohn/ML-foundations

Stay in Touch:

jonkrohn.com to sign up for email newsletter

 [linkedin.com/in/jonkrohn](https://www.linkedin.com/in/jonkrohn)

 [jonkrohn.com/youtube](https://www.youtube.com/jonkrohn)

 twitter.com/JonKrohnLearns



The Pomodoro Technique

Rounds of:

- 25 minutes of work
- with 5 minute breaks

Questions best handled at breaks, so save questions until then.

When people ask questions that have already been answered, do me a favor and let them know, politely providing response if appropriate.

Except during breaks, I recommend attending to this lecture only as topics are not discrete: Later material builds on earlier material.

POLL

Where are you?

- The Americas
- Europe / Middle East / Africa
- Asia-Pacific
- Extra-Terrestrial Space

POLL

What are you?

- Developer / Engineer
- Scientist / Analyst / Statistician / Mathematician
- Combination of the Above
- Other

POLL

What is your level of familiarity with Linear Algebra?

- Little to no exposure
- Some understanding of the theory
- Deep understanding of the theory
- Deep understanding of the theory and experience applying linear algebra operations with code

POLL

What is your level of familiarity with Machine Learning?

- Little to no exposure, or exposure to theory only
- Experience applying machine learning with code
- Experience applying machine learning with code and some understanding of the underlying theory
- Experience applying machine learning with code and strong understanding of the underlying theory

ML Foundations Series

Linear Algebra II builds upon and is foundational for:

1. Intro to Linear Algebra
2. **Linear Algebra II: Matrix Operations**
3. Calculus I: Limits & Derivatives
4. Calculus II: Partial Derivatives & Integrals
5. Probability & Information Theory
6. Intro to Statistics
7. Algorithms & Data Structures
8. Optimization

Linear Algebra II: Matrix Operations

1. Review of Introductory Linear Algebra
2. Eigendecomposition
3. Matrix Operations for Machine Learning

Linear Algebra II: Matrix Operations

1. **Review of Introductory Linear Algebra**
2. Eigendecomposition
3. Matrix Operations for Machine Learning

Segment 1: Review of Matrix Properties

- Modern Linear Algebra Applications
- Tensors, Vectors, and Norms
- Matrix Multiplication
- Matrix Inversion
- Identity, Diagonal and Orthogonal Matrices

What Linear Algebra Is

“Solving for unknowns within system of linear equations”


Consider the following example:

- Sheriff has 180 km/h car
- Bank robber has 150 km/h car and five-minute head start
- How long does it take the sheriff to catch the robber?
- What distance will they have traveled at that point?
- (For simplicity, let's ignore acceleration, traffic, etc.)

What Linear Algebra Is

Problem could be solved graphically with a plot:

(Note that: $150 \text{ km/h} = 2.5 \text{ km/min}$ $180 \text{ km/h} = 3 \text{ km/min}$)



What Linear Algebra Is

Alternatively, problem can be solved *algebraically*:

Equation 1: $d = 2.5t$

Equation 2: $d = 3(t - 5)$

$$2.5t = 3(t - 5)$$

$$2.5t = 3t - 15$$

$$2.5t - 3t = -15$$

$$-0.5t = -15$$

$$t = -15 / -0.5 = 30 \text{ min}$$



$$d = 2.5t = 2.5(30) = 75 \text{ km}$$

$$d = 3(t - 5) = 3(30 - 5) = 3(25) = 75 \text{ km}$$

What Linear Algebra Is

No solution if sheriff's car is same speed as bank robber's.

Infinite solutions if same speed *and* same starting time.



These are the only three options in linear algebra: one, no, or infinite solutions.

It is impossible for lines to cross multiple times.

What Linear Algebra Is

In a given system of equations:

- Could be *many* equations
- Could be *many* unknowns in each equation

A handwritten linear regression equation with several annotations. The equation is $y = a + bx_1 + cx_2 + \dots + mx_m$. Annotations include: a blue arrow pointing from "house price" to y ; an orange arrow pointing from "distance to school" to x_1 ; a green arrow pointing from "y-intercept" to a ; a brown arrow pointing from "number of bedrooms" to x_2 ; and a grey arrow pointing from "there could be m features (many!)" to the ellipsis and x_m .

$$y = a + bx_1 + cx_2 + \dots + mx_m$$

house price

distance to school


there could be m features (many!)

"y-intercept"

number of bedrooms

$$y = a + b x_1 + c x_2 + \dots + m x_m$$

$$\left[\begin{array}{c|c} y_1 & a + b x_{1,1} + c x_{1,2} + \dots + m x_{1,m} \\ y_2 & a + b x_{2,1} + c x_{2,2} + \dots + m x_{2,m} \\ \vdots & \vdots \\ y_n & a + b x_{n,1} + c x_{n,2} + \dots + m x_{n,m} \end{array} \right]$$


 For any house i in the dataset,
 y_i = price and $x_{i,1}$ to $x_{i,m}$ are its features.
 We solve for parameters a, b, c to m

Modern Linear Algebra Applications

- Solving for unknowns in ML algos, including deep learning
- Reducing dimensionality (e.g., **principal component analysis**)
- Ranking results (e.g., with **eigenvector**, including in Google PageRank algorithm; *see Saaty and Hu, 1998*)
- Recommenders (e.g., **singular value decomposition, SVD**)
- Natural language processing (e.g., **SVD**, matrix factorization)
 - Topic modeling
 - Semantic analysis

Tensors

“ML generalization of vectors and matrices to any number of dimensions”

scalar


x

vector

$[x_1 \ x_2 \ x_3]$

matrix

$\begin{bmatrix} x_{1,1} & x_{1,2} \\ x_{2,1} & x_{2,2} \end{bmatrix}$



3-tensor

Dimensions	Mathematical Name	Description
0	scalar	magnitude only
1	vector	array
2	matrix	flat table, e.g., square
3	3-tensor	3D table, e.g., cube
n	n -tensor	higher dimensional

Vector Transposition

$$\begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix}^T = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

row vector

column
vector


shape is (1, 3)

(3, 1)

Hands-on code demo: [2-linear-algebra-ii.ipynb](#)

Norms

Vectors represent a magnitude and direction from origin:



Norms are functions that quantify vector magnitude:

- In ML, L^2 and L^1 norms are common, e.g., to avoid overfitting

L^2 Norm

- Described by:


$$\| \mathbf{x} \|_2 = \sqrt{\sum_i x_i^2}$$

- Measures simple (Euclidean) distance from origin
- Most common norm in machine learning
 - Instead of $\|\mathbf{x}\|_2$, it can be denoted as $\|\mathbf{x}\|$

Hands-on code demo


Unit Vectors

- Special case of vector where its length is equal to one
- Technically, \mathbf{x} is a unit vector with “unit norm”, i.e.: $\|\mathbf{x}\| = 1$




Basis Vectors

- Can be scaled to represent *any* vector in a given vector space
- Typically use unit vectors along axes of vector space (shown)



Orthogonal Vectors

- \mathbf{x} and \mathbf{y} are orthogonal vectors if $\mathbf{x}^T \mathbf{y} = 0$
- Are at 90° angle to each other (assuming non-zero norms)
- n -dimensional space has max n mutually orthogonal vectors (again, assuming non-zero norms)
- **Orthonormal** vectors are orthogonal *and* all have unit norm
 - Basis vectors are an example



Matrix Transposition

Flip of axes over **main diagonal** such that:

$$(X^T)_{i,j} = X_{j,i}$$

$$\begin{bmatrix} x_{1,1} & x_{1,2} \\ x_{2,1} & x_{2,2} \\ x_{3,1} & x_{3,2} \end{bmatrix}^T = \begin{bmatrix} x_{1,1} & x_{2,1} & x_{3,1} \\ x_{1,2} & x_{2,2} & x_{3,2} \end{bmatrix}$$

Hands-on code demo

Symmetric Matrices

Special matrix case with following properties:

- Square
- $\mathbf{X}^T = \mathbf{X}$

$$\begin{bmatrix} 0 & 1 & 2 \\ 1 & 7 & 8 \\ 2 & 8 & 9 \end{bmatrix}$$

Matrix Multiplication

$$\begin{matrix} m \\ \left[\begin{array}{c} C \end{array} \right] \\ p \end{matrix} = \begin{matrix} m \\ \left[\begin{array}{c} A \end{array} \right] \\ n \end{matrix} \begin{matrix} n \\ \left[\begin{array}{c} B \end{array} \right] \\ p \end{matrix}$$

$$C_{i,k} = \sum_j A_{i,j} B_{j,k}$$

Matrix Multiplication (with a Vector)

$$\begin{bmatrix} 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 3 \cdot 1 + 4 \cdot 2 \\ 5 \cdot 1 + 6 \cdot 2 \\ 7 \cdot 1 + 8 \cdot 2 \end{bmatrix} = \begin{bmatrix} 3 + 8 \\ 5 + 12 \\ 7 + 16 \end{bmatrix} = \begin{bmatrix} 11 \\ 17 \\ 23 \end{bmatrix}$$

Hands-on code demo

(Matrix-by-)Matrix Multiplication

$$\begin{bmatrix} 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} \begin{bmatrix} 1 & 9 \\ 2 & 0 \end{bmatrix} = \begin{bmatrix} 3 \cdot 1 + 4 \cdot 2 & 3 \cdot 9 + 4 \cdot 0 \\ 5 \cdot 1 + 6 \cdot 2 & 5 \cdot 9 + 6 \cdot 0 \\ 7 \cdot 1 + 8 \cdot 2 & 7 \cdot 9 + 8 \cdot 0 \end{bmatrix} = \begin{bmatrix} 11 & 27 \\ 17 & 45 \\ 23 & 63 \end{bmatrix}$$

Hands-on code demo

$$y = a + b x_1 + c x_2 + \dots + m x_m$$

$$\left[\begin{array}{c|c} y_1 & a + b x_{1,1} + c x_{1,2} + \dots + m x_{1,m} \\ y_2 & a + b x_{2,1} + c x_{2,2} + \dots + m x_{2,m} \\ \vdots & \vdots \\ y_n & a + b x_{n,1} + c x_{n,2} + \dots + m x_{n,m} \end{array} \right]$$

Strictly speaking, x extends rightward to $m-1$ not m because of the presence of a on the far left.

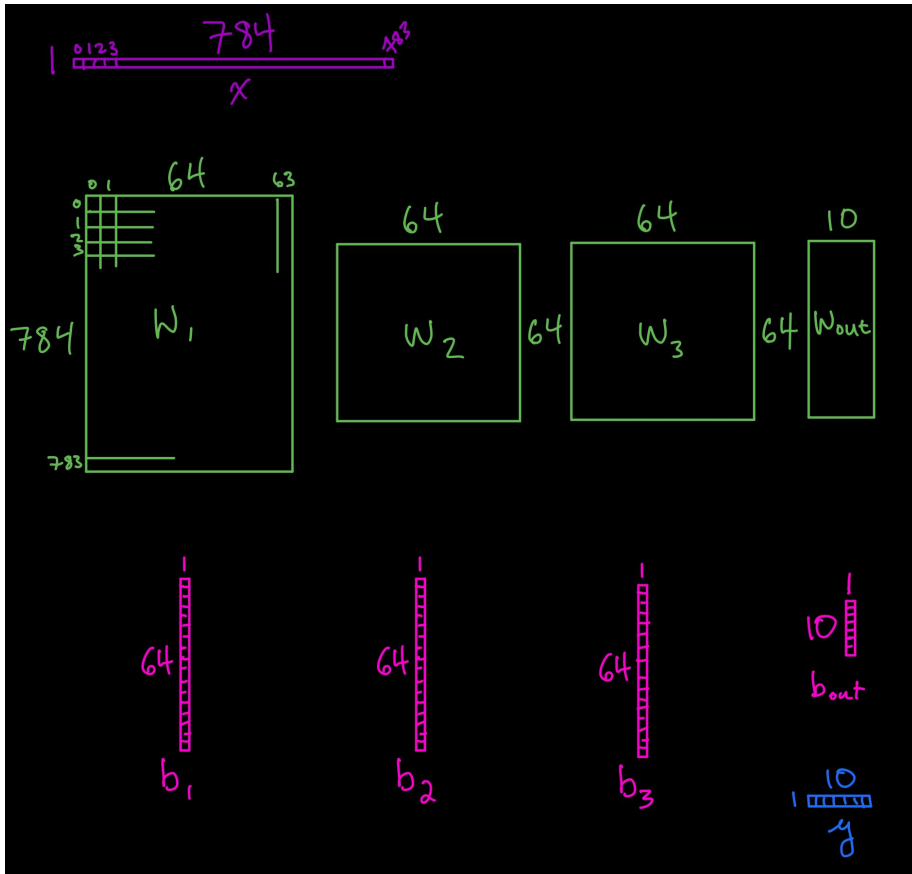
For any house i in the dataset,
 y_i = price and $x_{i,1}$ to $x_{i,m}$ are its features.
 We solve for parameters a, b, c to m

Matrix Multiplication (in Regression)

$$\begin{array}{c} \text{\textit{n}} \\ \text{cases} \\ \text{tall} \end{array} \left\{ \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \right\} = \begin{bmatrix} | & x_{1,1} & x_{1,2} & \cdots & x_{1,m} \\ | & x_{2,1} & x_{2,2} & \cdots & x_{2,m} \\ \vdots & \vdots & \vdots & & \vdots \\ | & x_{n,1} & x_{n,2} & \cdots & x_{n,m} \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ \vdots \\ m \end{bmatrix}$$

$\underbrace{\hspace{15em}}_{m \text{ features wide}}$

Matrix Multiplication (in Deep Learning)



See:

- [artificial-neurons.ipynb](#)
- [jonkrohn.com/deepTF1](#)
- [jonkrohn.com/convTF1](#)
- [jonkrohn.com/convTF2](#)
- [jonkrohn.com/deepPT](#)

Identity Matrices

Symmetric matrix where:

- Every element along main diagonal is 1
- All other elements are 0
- Notation: \mathbf{I}_n where n = height (or width)
- n -length vector unchanged if multiplied by \mathbf{I}_n

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

\mathcal{I}_4

Matrix Inversion

- Clever, convenient approach for solving linear equations
- An alternative to manually solving with elimination or addition
- **Matrix inverse** of X is denoted as X^{-1}
 - Satisfies: $X^{-1}X = XX^{-1} = I_n$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

I_4

Matrix Inversion

$$\begin{matrix} n \\ \text{cases} \\ \text{tall} \end{matrix} \left\{ \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \right\} = \underbrace{\begin{bmatrix} | & x_{1,1} & x_{1,2} & \dots & x_{1,m} \\ | & x_{2,1} & x_{2,2} & \dots & x_{2,m} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ | & x_{n,1} & x_{n,2} & \dots & x_{n,m} \end{bmatrix}}_{m \text{ features wide}} \begin{bmatrix} a \\ b \\ c \\ \vdots \\ m \end{bmatrix}$$

The regression formula can be represented as:

$$\mathbf{y} = \mathbf{X}\mathbf{w} \quad (\mathbf{w} \text{ is the vector of weights } b \text{ through } m)$$

Matrix Inversion

In the equation $\mathbf{y} = \mathbf{X}\mathbf{w}$:

- We know the outcomes \mathbf{y} , which could be house prices
- We know the features \mathbf{X} , which are predictors like bedroom count
- Vector \mathbf{w} contains the unknowns, the model's learnable parameters

Assuming \mathbf{X}^{-1} exists, matrix inversion can solve for \mathbf{w} :

$$\mathbf{X}\mathbf{w} = \mathbf{y}$$

$$\mathbf{X}^{-1}\mathbf{X}\mathbf{w} = \mathbf{X}^{-1}\mathbf{y}$$

$$\mathbf{I}_n \mathbf{w} = \mathbf{X}^{-1}\mathbf{y}$$

$$\mathbf{w} = \mathbf{X}^{-1}\mathbf{y}$$

Matrix Inversion

$$\begin{cases} 4b + 2c = 4 \\ -5b - 3c = -7 \end{cases}$$

$$X = \begin{bmatrix} x_{1,1} & x_{1,2} \\ x_{2,1} & x_{2,2} \end{bmatrix} = \begin{bmatrix} 4 & 2 \\ -5 & -3 \end{bmatrix} \quad y = \begin{bmatrix} 4 \\ -7 \end{bmatrix}$$


$$w = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} b \\ c \end{bmatrix} = X^{-1}y$$

Hands-on code demo

Matrix Inversion

Nifty trick, but can only be calculated if:

- Matrix is square: $n_{\text{row}} = n_{\text{col}}$ (i.e., “vector span” = “matrix range”)
 - Avoids **overdetermination**: $n_{\text{row}} (\# \text{ of equations}) > n_{\text{col}} (\# \text{ of dims})$
 - Avoids **underdetermination**: $n_{\text{row}} < n_{\text{col}}$
- Matrix isn't “singular”, i.e.: all columns are linearly independent
 - E.g., if a column is $[1, 2]$, another can't be $[2, 4]$ or also be $[1, 2]$**



Note that solving for unknowns may still be possible by other means if matrix can't be inverted...

Diagonal Matrices

- Nonzero elements along main diagonal; zeros everywhere else
- Identity matrix is an example
- If square, denoted as $\text{diag}(\mathbf{x})$ where \mathbf{x} is vector of main-diagonal elements
- Computationally efficient:
 - Multiplication: $\text{diag}(\mathbf{x})\mathbf{y} = \mathbf{x} \odot \mathbf{y}$
 - Inversion: $\text{diag}(\mathbf{x})^{-1} = \text{diag}[1/\mathbf{x}_1, \dots, 1/\mathbf{x}_n]^T$
 - Can't divide by zero so \mathbf{x} can't include zero
- Can be non-square and computation still efficient:
 - If $h > w$, simply add zeros to product
 - If $w > h$, remove elements from product

A hand-drawn 4x4 identity matrix is shown, enclosed in large square brackets. The matrix has ones on the main diagonal and zeros elsewhere. The elements are represented as follows: Row 1: 1, 0, 0, 0; Row 2: 0, 1, 0, 0; Row 3: 0, 0, 1, 0; Row 4: 0, 0, 0, 1. Below the matrix, the label I_4 is written in a cursive script.

Orthogonal Matrices

Recall orthonormal vectors from earlier:


In orthogonal matrices, orthonormal vectors:

- Make up all rows
- Make up all columns

This means: $\mathbf{A}^T \mathbf{A} = \mathbf{A} \mathbf{A}^T = \mathbf{I}$

Which also means: $\mathbf{A}^T = \mathbf{A}^{-1} \mathbf{I} = \mathbf{A}^{-1}$

Calculating \mathbf{A}^T is cheap, therefore so is calculating \mathbf{A}^{-1}



Linear Algebra II: Matrix Operations

1. Review of Introductory Linear Algebra
2. **Eigendecomposition**
3. Matrix Operations for Machine Learning

Segment 2: Eigendecomposition

- Applying Matrices
- Affine Transformations
- Eigenvectors
- Eigenvalues
- Matrix Determinants
- Matrix Decomposition
- Applications of Eigendecomposition

Matrix-Application Exercises

Using pen(cil) and paper:

1. Apply the identity matrix \mathbf{I}_3 to the vector \mathbf{u} .
2. Apply the matrix \mathbf{B} to the vector \mathbf{u} .
3. Concatenate vector \mathbf{u} with vector \mathbf{u}_2 to form a matrix \mathbf{U} , then apply the matrix \mathbf{B} to the matrix \mathbf{U} .

$$\mathbf{u} = \begin{bmatrix} 2 \\ 5 \\ -3 \end{bmatrix} \quad \mathbf{u}_2 = \begin{bmatrix} 0 \\ -4 \\ 6 \end{bmatrix}$$

$$\mathbf{B} = \begin{bmatrix} 2 & 0 & -1 \\ -2 & 3 & 1 \\ 0 & 4 & -1 \end{bmatrix}$$

Solutions


$$I_3 \mu = \begin{bmatrix} 2 \\ 5 \\ -3 \end{bmatrix}$$

$$B_{\mu} = \begin{bmatrix} 4 & +0 & +3 \\ -4 & +15 & -3 \\ 0 & 20 & 3 \end{bmatrix} = \begin{bmatrix} 7 \\ 8 \\ 23 \end{bmatrix}$$

$$B_{\mu_2} = \begin{bmatrix} 0 & +0 & -6 \\ 0 & -12 & 6 \\ 0 & -16 & -6 \end{bmatrix} = \begin{bmatrix} -6 \\ -6 \\ -22 \end{bmatrix}$$

$$BU = \begin{bmatrix} 7 & -6 \\ 8 & -6 \\ 23 & -22 \end{bmatrix}$$


Eigenvectors




Flipping matrix applied



Red vector and **blue vector**
are **eigenvectors** for the
flipping matrix.



Eigenvectors




Shearing matrix applied




Red vector knocked off span

Blue vector isn't -- it maintains its direction -- so it is an **eigenvector** for the shearing matrix.




Eigenvalues




In this case, **eigenvector** retains exact length, so its **eigenvalue** = 1.

If **eigenvector** were to double in length, its **eigenvalue** = 2; if it halves, **eigenvalue** = 0.5.




Eigenvalues



Eigenvalues can also have a negative sign, e.g., a new **shearing-and-flipping matrix** has the same **eigenvector** as shearing-only matrix but its **eigenvalue** = -1.



If **eigenvector** were to double in length while exactly reversing direction, **eigenvalue** would be -2.

Hands-on code demo



Matrix Determinants

- Map square matrix to scalar
- Enable us to determine whether matrix can be inverted
- For matrix X , denoted as $\det(X)$
- If $\det(X) = 0$:
 - Matrix X^{-1} can't be computed because: X^{-1} has $1/\det(X) = 1/0$
 - Matrix X is singular: It contains linearly-dependent columns
- $\det(X)$ easiest to calculate for 2x2 matrix...



Determinant of 2x2 Matrix

$$X = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

$$|X| = ad - bc$$

$$X = \begin{bmatrix} 4 & 2 \\ -5 & -3 \end{bmatrix}$$

$$\begin{aligned} |X| &= 4(-3) - 2(-5) \\ &= -12 + 10 \\ &= -2 \end{aligned}$$

Hands-on code demo

Determinant of 2x2 Matrix

$$X = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

$$|X| = ad - bc$$

$$n = \begin{bmatrix} -4 & 1 \\ -8 & 2 \end{bmatrix}$$

$$\begin{aligned} |n| &= -4 \cdot 2 - 1(-8) \\ &= -8 + 8 \\ &= 0 \end{aligned}$$

Hands-on code demo

Generalizing Determinants: Recursion

$$X = \begin{bmatrix} x_{1,1} & x_{1,2} & x_{1,3} & x_{1,4} & x_{1,5} \\ x_{2,1} & x_{2,2} & x_{2,3} & x_{2,4} & x_{2,5} \\ x_{3,1} & x_{3,2} & x_{3,3} & x_{3,4} & x_{3,5} \\ x_{4,1} & x_{4,2} & x_{4,3} & x_{4,4} & x_{4,5} \\ x_{5,1} & x_{5,2} & x_{5,3} & x_{5,4} & x_{5,5} \end{bmatrix}$$

} 5 rows
means
4 rounds
of
recursion

$$|X| = x_{1,1} \det(x_{1,1}) - x_{1,2} \det(x_{1,2}) + x_{1,3} \det(x_{1,3}) \\ - x_{1,4} \det(x_{1,4}) + x_{1,5} \det(x_{1,5})$$

alternating +/-

Generalizing Determinants: Recursion

$$\begin{aligned} X &= \begin{bmatrix} 1 & 2 & 4 \\ 2 & -1 & 3 \\ 0 & 5 & 1 \end{bmatrix} \left. \vphantom{\begin{bmatrix} 1 & 2 & 4 \\ 2 & -1 & 3 \\ 0 & 5 & 1 \end{bmatrix}} \right\} \begin{array}{l} 3 \text{ rows means} \\ 2 \text{ rounds of recursion} \end{array} \\ |X| &= x_{1,1} \det(X_{1,1}) - x_{1,2} \det(X_{1,2}) + x_{1,3} \det(X_{1,3}) \\ &= 1 \begin{vmatrix} -1 & 3 \\ 5 & 1 \end{vmatrix} - 2 \begin{vmatrix} 2 & 3 \\ 0 & 1 \end{vmatrix} + 4 \begin{vmatrix} 2 & -1 \\ 0 & 5 \end{vmatrix} \\ &= 1(-1 \cdot 1 - 3 \cdot 5) - 2(2 \cdot 1 - 3 \cdot 0) + 4(2 \cdot 5 - (-1)(0)) \\ &= 1(-1 - 15) - 2(2 - 0) + 4(10 - 0) \\ &= -16 - 4 + 40 \\ &= 20 \end{aligned}$$

Hands-on code demo

Exercises

Using pencil and paper, calculate the determinant of the matrices below. Indicate which have an inverse and which don't:

1. $\begin{bmatrix} 25 & 2 \\ 3 & 4 \end{bmatrix}$

2. $\begin{bmatrix} -2 & 0 \\ 0 & -2 \end{bmatrix}$

3. $\begin{bmatrix} 2 & 1 & -3 \\ 4 & -5 & 2 \\ 0 & -1 & 3 \end{bmatrix}$

Solutions

1. 94; has inverse
2. 4; has inverse
3. -26; has inverse

Determinants & Eigenvalues

$\det(\mathbf{X})$ = product of all eigenvalues of \mathbf{X}

Hands-on code demo

$|\det(\mathbf{X})|$ quantifies volume change as a result of applying \mathbf{X} :

- If $\det(\mathbf{X}) = 0$, then \mathbf{X} collapses space completely in at least one dimension, thereby eliminating all volume
- If $0 < |\det(\mathbf{X})| < 1$, then \mathbf{X} contracts volume to some extent
- If $|\det(\mathbf{X})| = 1$, then \mathbf{X} preserves volume exactly
- If $|\det(\mathbf{X})| > 1$, then \mathbf{X} expands volume

Hands-on code demo

Eigendecomposition




$$A = V\Lambda V^{-1}$$

The decomposition of a matrix into eigenvectors and eigenvalues reveals characteristics of the matrix, e.g.:

- Matrix is singular if and only if any of its eigenvalues are zero
- Under specific conditions (see §2.7 of Goodfellow et al., 2016), can optimize quadratic expressions:
 - Maximum of $f(\mathbf{x})$ = largest eigenvalue
 - Minimum of $f(\mathbf{x})$ = smallest eigenvalue

Eigendecomposition Examples

2D geometric transformation	Scaling (equal)	Scaling (unequal)	Horizontal shear	Vertical shear
2x2 Matrix	$\begin{bmatrix} k & 0 \\ 0 & k \end{bmatrix}$	$\begin{bmatrix} k_1 & 0 \\ 0 & k_2 \end{bmatrix}$	$\begin{bmatrix} 1 & k \\ 0 & 1 \end{bmatrix}$	$\begin{bmatrix} 1 & 0 \\ k & 1 \end{bmatrix}$
Eigenvalues	$\lambda_1 = \lambda_2 = k$	$\lambda_1 = k_1$ and $\lambda_2 = k_2$	$\lambda_1 = \lambda_2 = 1$	$\lambda_1 = \lambda_2 = 1$
Example eigenvectors	non-zero	$\mathbf{v}_1 = [1, 0]$ and $\mathbf{v}_2 = [0, 1]$	$\mathbf{v}_1 = [1, 0]$	$\mathbf{v}_1 = [0, 1]$



Eigendecomposition Applications


Matrix is of type:	If all its eigenvalues are:
Positive definite	>0
Positive semidefinite	≥ 0
Negative definite	<0
Negative semidefinite	≤ 0

Applying a matrix of a particular type to some vector \mathbf{x} can have a characteristic impact (again, see §2.7 of Goodfellow et al., 2016):

- E.g., semidefinite matrices collapse tensors along 1+ dimensions

Eigendecomposition Applications

- Eigenvectors, as underlying characteristics of a dataset, can be recombined into any members of the dataset, e.g.:
 - Eigenfaces (**shown**)
 - Eigenvoices
 - Eigenfrequencies (of vibrations)
- Quantum mechanics:
 - Molecular orbitals
 - Schrödinger wave equation
- Reproduction number R_0 in epidemiology
- Calculating determinants (*already covered*)
- SVD & Moore-Penrose pseudoinverse (*next*)
- Principal component analysis (*coming up*)



Linear Algebra II: Matrix Operations

1. Review of Introductory Linear Algebra
2. Eigendecomposition
3. **Matrix Operations for Machine Learning**

Segment 3: Matrix Operations for ML

- Singular Value Decomposition (SVD)
- The Moore-Penrose Pseudoinverse
- The Trace Operator
- Principal Component Analysis (PCA)
- Resources for Further Study of Linear Algebra

Singular Value Decomposition

- Unlike eigendecomposition, which is applicable to square matrices only, SVD is applicable to *any* real-valued matrix
- Decomposes matrix into:
 - **Singular vectors** (analogous to eigenvectors)
 - **Singular values** (analogous to eigenvalues)
- For some matrix A , its SVD is $A = UDV^T$


Hands-on code demo

$$\begin{matrix} m \\ \left[\begin{array}{c} A \end{array} \right] \\ n \end{matrix} = \begin{matrix} m \\ \left[\begin{array}{c} \text{left-singular vector} \\ U \end{array} \right] \\ m \end{matrix} \begin{matrix} \left[\begin{array}{c} \text{singular values} \\ D \end{array} \right] \\ n \end{matrix} \begin{matrix} \left[\begin{array}{c} \text{R-singular vec} \\ V^T \end{array} \right] \\ n \end{matrix}$$

Matrix Inversion *Revisited*

Nifty trick, but can only be calculated if:

- Matrix is square: $n_{\text{row}} = n_{\text{col}}$ (i.e., “vector span” = “matrix range”)
 - Avoids **overdetermination**: $n_{\text{row}} (\# \text{ of equations}) > n_{\text{col}} (\# \text{ of dims})$
 - Avoids **underdetermination**: $n_{\text{row}} < n_{\text{col}}$
- Matrix isn't “singular”, i.e.: all columns are linearly independent
 - **E.g., if a column is $[1, 2]$, another can't be $[2, 4]$ or also be $[1, 2]$**



Note that solving for unknowns may still be possible by other means if matrix can't be inverted.

Moore-Penrose Pseudoinverse

- Such an “other mean” when matrix itself can’t be inverted
- Instead, invert the singular values of the matrix

For some matrix A , its pseudoinverse A^+ can be calculated by:

$$A^+ = VD^+U^T$$

Where:

- U , D , and V are SVD of A
- $D^+ = (D \text{ with reciprocal of all-non zero elements})^T$

Hands-on code demo

Moore-Penrose Pseudoinverse

A^+ is mega useful because non-square matrices are common in ML:

$$y = a + b x_1 + c x_2 + \dots + m x_m$$

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \begin{bmatrix} a + b x_{1,1} + c x_{1,2} + \dots + m x_{1,m} \\ a + b x_{2,1} + c x_{2,2} + \dots + m x_{2,m} \\ \vdots \\ a + b x_{n,1} + c x_{n,2} + \dots + m x_{n,m} \end{bmatrix}$$

↖ For any house i in the dataset,
 y_i = price and $x_{i,1}$ to $x_{i,m}$ are its features.
We solve for parameters a, b, c to m

Matrix Inversion *Revisited*

$$\begin{array}{c} n \\ \text{cases} \\ \text{tall} \end{array} \left\{ \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \right\} = \underbrace{\begin{bmatrix} | & x_{1,1} & x_{1,2} & \cdots & x_{1,m} \\ | & x_{2,1} & x_{2,2} & \cdots & x_{2,m} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ | & x_{n,1} & x_{n,2} & \cdots & x_{n,m} \end{bmatrix}}_{m \text{ features wide}} \begin{bmatrix} a \\ b \\ c \\ \vdots \\ m \end{bmatrix}$$

The regression formula can be represented as:

$$\mathbf{y} = \mathbf{X}\mathbf{w} \quad (\mathbf{w} \text{ is the vector of weights } a \text{ through } m)$$

Matrix Inversion *Revisited*

In the equation $y = Xw$:

- We know the outcomes y , which could be house prices
- We know the features X , which are predictors like bedroom count
- Vector w contains the unknowns, the model's learnable parameters

Assuming X^{-1} exists, matrix inversion can solve for w :

$$Xw = y$$

$$X^{-1}Xw = X^{-1}y$$

$$I_n w = X^{-1}y$$

$$w = X^{-1}y$$

Matrix Inversion Revisited

$$\begin{cases} 4b + 2c = 4 \\ -5b - 3c = -7 \end{cases}$$

$$X = \begin{bmatrix} x_{1,1} & x_{1,2} \\ x_{2,1} & x_{2,2} \end{bmatrix} = \begin{bmatrix} 4 & 2 \\ -5 & -3 \end{bmatrix} \quad y = \begin{bmatrix} 4 \\ -7 \end{bmatrix}$$

$$w = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} b \\ c \end{bmatrix} = X^{-1}y$$



Moore-Penrose Pseudoinverse

- Would be unusual to have exactly as many cases (n) as features (m)
- With pseudoinverse \mathbf{X}^+ , we can now estimate model weights \mathbf{w} if $n \neq m$:

$$\mathbf{w} = \mathbf{X}^+ \mathbf{y}$$

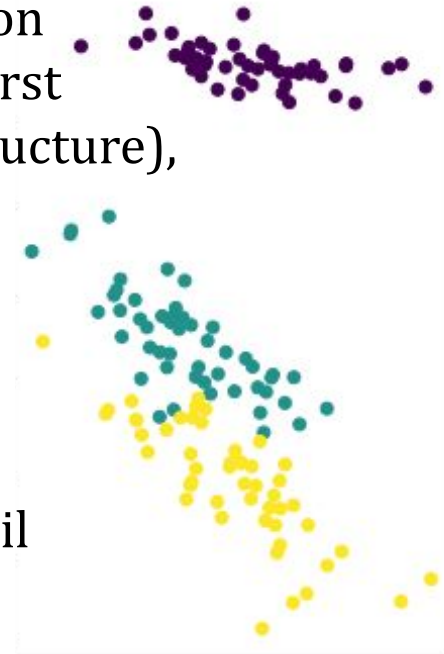
- If \mathbf{X} is **overdetermined** ($n > m$), \mathbf{X}^+ provides $\mathbf{X}\mathbf{y}$ as close to \mathbf{w} as possible (in terms of Euclidean distance, specifically $\|\mathbf{X}\mathbf{y} - \mathbf{w}\|_2$)
- If \mathbf{X} is **underdetermined** ($n < m$), \mathbf{X}^+ provides the $\mathbf{w} = \mathbf{X}^+ \mathbf{y}$ solution that has the smallest Euclidean norm $\|\mathbf{x}\|_2$ from all the possible solutions

Hands-on code demo



Principal Component Analysis

- Simple machine learning algorithm
- **Unsupervised**: enables identification of structure in unlabeled data
- Like eigendecomposition and SVD, enables lossy compression
 - To minimize both loss of precision and data footprint, first **principal component** contains most variance (data structure), second PC contains next most, and so on
- Involves many linear algebra concepts already covered, e.g.:
 - Norms
 - Orthogonal and identity matrices
 - Trace operator
 - See Goodfellow et al. (2016) §2.12 for five pages of detail



Resources for Further Study

- **Basic algebra:**
 - Khan Academy
 - 3Blue1Brown on YouTube
- **Linear algebra:**
 - 3Blue1Brown again
 - Ch. 2 of Goodfellow et al. (2016) *Deep Learning* ([free](#))
 - Ch. 2 of Deisenroth et al. (2020) [*Mathematics for ML*](#)
 - Sheldon Axler's (2015) [*Linear Algebra Done Right*](#)
 - Gilbert Strang: [Linear Algebra course](#) via MIT Open Courseware
 - *Linear Algebra and Its Applications* book
- **Next steps in the *ML Foundations* series:**
 - Calculus I: Limits & Derivatives
 - Calculus II: Partial Derivatives & Integrals

Stay in Touch

jonkrohn.com to sign up for email newsletter



[linkedin.com/in/jonkrohn](https://www.linkedin.com/in/jonkrohn)



YouTube jonkrohn.com/youtube



twitter.com/JonKrohnLearns



POLL *with Multiple Answers Possible*

What follow-up topics interest you most?

- More Linear Algebra
- Calculus
- Probability / Statistics
- Computer Science (e.g., algorithms, data structures)
- Machine Learning Basics
- Advanced Machine Learning, incl. Deep Learning
- Something Else



NEBULA

PLACEHOLDER
FOR:

5-Minute Timer

PLACEHOLDER
FOR:

10-Minute Timer

PLACEHOLDER
FOR:

15-Minute Timer